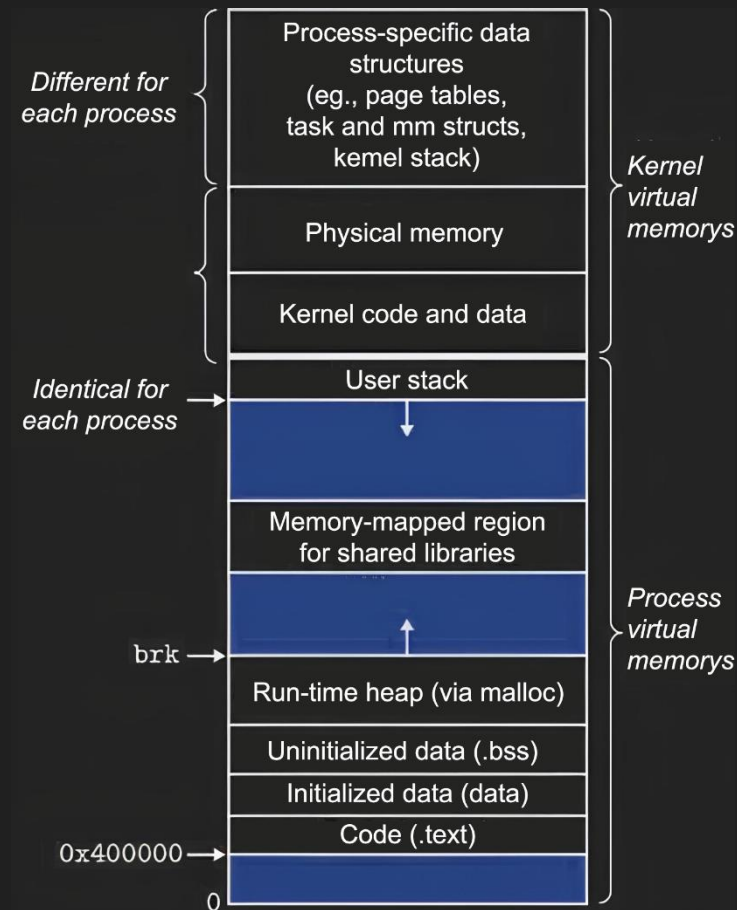


WEEK 1

Memory layout,
internal object representation

Virtual memory layout



Stack

- Memory allocation is "free" and "continuous"

```
int data[10]; -> subq $48, %rsp
```

- Usually limited in size, but can be increased
- Allocation (stack growth) can lead to a page fault, so there is a potential cost to it, but its 1-time only, since stack doesn't shrink

Heap

- An allocation here is time-consuming in general
- The performance can degrade as the application continues to run
- There are a lot of different allocators, each is designed around certain workload it needs to handle, so there is no silver bullet
- Allocations can easily take most of the running time if you don't pay attention
- Most allocators have really bad tail-latency values.

General performance tips

- time vs memory consumption tradeoff
- In general you want to preallocate as much memory as possible (in hft the entire memory pool is static or allocated on heap at the start)
- If you have to increase the default stack size, it usually means there is something wrong with your code.
- Consider what exactly happens when you use some std-container. Do not replace blindly `std::vector` with `std::array` just because its "faster" (its not).
- Don't mix memory management logic with the rest of your code. You can still have nice std-like interfaces and 0 allocations (home tasks).

Memory layout of simple types

Assume little-endian byte order when not specified in this course

```
char a = 1;  
int b = 17239;  
int c = -17239;  
const char *d = "text";  
char e[10];  
std::cin >> e; // "hello"  
double f = 2.0;  
int64_t g = 3;
```

```
(gdb) x/60b &g  
0x7fffffff108: g<0x03 0x00 0x00 0x00 0x00 0x00 0x00 0x00>  
0x7fffffff110: f<0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x40>  
0x7fffffff118: 0x10 0x12 0x40 0x00 0x00 0x00 e<0x68 0x65  
0x7fffffff120: 0x6c 0x6c 0x6f 0x00 0x00 0x00 0x00 0x00>  
0x7fffffff128: d<0x10 0x20 0x40 0x00 0x00 0x00 0x00 0x00>  
0x7fffffff130: 0x30 0xe2 0xff 0xff c<0xa9 0xbc 0xff 0xff>  
0x7fffffff138: b<0x57 0x43 0x00 0x00> 0x00 0x00 0x00 a<0x01>  
0x7fffffff140: 0x00 0x00 0x00 0x00
```

Memory layout and padding

```
struct Item {  
    int64_t a;  
    int b;  
    char c;  
};  
...  
std::vector<Item> data{Item{1, 2, 'a'}, Item{-1, 64, 'b'}};  
std::cout << sizeof(Item) << " " << alignof(Item) << "\n"; // 16 8
```

```
(gdb) x/8wx &data[0]  
0x416eb0:    0x00000001    0x00000000    0x00000002    0x00000061  
0x416ec0:    0xffffffff    0xffffffff    0x00000040    0x00000062
```

Disable padding

```
#pragma pack(push, 1)
struct Item {
    int64_t a;
    int b;
    char c;
};
#pragma pack(pop) // size = 13, alignof = 1
```

Not recommended, since it slows down any memory access

Memory layout of complex classes

```
struct Item {  
    std::vector<int> a;  
    std::string b;  
};
```

```
Item a{std::vector<int>{1, 2, 3}, "hello"};  
std::cout << sizeof(a.a) << " " << sizeof(a.b) << "\n"; // 24 32
```

```
(gdb) x/60bx &a
```

0x7fffffff108:	0xb0	0x6e	0x41	0x00	0x00	0x00	0x00	0x00
0x7fffffff110:	0xbc	0x6e	0x41	0x00	0x00	0x00	0x00	0x00
0x7fffffff118:	0xbc	0x6e	0x41	0x00	0x00	0x00	0x00	0x00
0x7fffffff120:	0x30	0xe1	0xff	0xff	0xff	0x7f	0x00	0x00
0x7fffffff128:	0x05	0x00	0x00	0x00	0x00	0x00	0x00	0x00
0x7fffffff130:	0x68	0x65	0x6c	0x6c	0x6f	0x00	0x00	0x00
0x7fffffff138:	0x00	0x00	0x00	0x00	0x00	0x00	0x00	0x00
0x7fffffff140:	0x00	0x00	0x00	0x00				

```
(gdb) p &a.a[0]
```

```
$1 = (int *) 0x416eb0
```

Functions and classes (without inheritance)

- Functions (i.e the instructions) are located in text section of a file and so those are stored separately from the data.
- Method is basically a function with `this` as it's first argument.
- So the layout of any class is defined by the layout of it's members.
- `public` , `private` sections do not matter.
- templates is just an advance codegen mechanism, doesn't affect the layout

Inheritance

```
struct A {  
    int a;  
    explicit A(int _a): a(_a) {}  
    virtual int f() { return 5; }  
};  
struct B : public A {  
    int b;  
    virtual int f() { return 7; }  
    B(int a, int b): A(a), b(b) {}  
};
```

```
A a{1};  
B b{2, 3};
```

```
(gdb) x/40bx &b  
0x7fffffffef120: 0x48      0x20      0x40      0x00      0x00      0x00      0x00      0x00  
0x7fffffffef128: 0x02      0x00      0x00      0x00      0x03      0x00      0x00      0x00  
0x7fffffffef130: 0x18      0x20      0x40      0x00      0x00      0x00      0x00      0x00  
0x7fffffffef138: 0x01      0x00      0x00      0x00      0x00      0x00      0x00      0x00  
0x7fffffffef140: 0x00      0x00      0x00      0x00      0x00      0x00      0x00      0x00
```

Inheritance

```
B b{2, 3};  
A *a = &b;  
int c = a->f();
```

```
(gdb) x/24bx &a
```

0x7fffffff128: 0x30	0xe1	0xff	0xff	0xff	0x7f	0x00	0x00
0x7fffffff130: 0x18	0x20	0x40	0x00	0x00	0x00	0x00	0x00
0x7fffffff138: 0x02	0x00	0x00	0x00	0x03	0x00	0x00	0x00

```
(gdb) x/24bx 0x402018
```

0x402018 <_ZTV1B+16>:	0x60	0x12	0x40	0x00	0x00	0x00	0x00	0x00
0x402020 <_ZTS1B>:	0x31	0x42	0x00	0x31	0x41	0x00	0x00	0x00
0x402028 <_ZTI1A>:	0x50	0x3d	0x40	0x00	0x00	0x00	0x00	0x00

```
(gdb) x/24bx 0x401260
```

0x401260 <B::f()>:	0x55	0x48	0x89	0xe5	0x48	0x89	0x7d	0xf8
0x401268 <B::f()+8>:	0xb8	0x07	0x00	0x00	0x00	0x5d	0xc3	0x90
0x401270 <A::f()>:	0x55	0x48	0x89	0xe5	0x48	0x89	0x7d	0xf8

Inheritance

```
(gdb) x/15i 0x401260
0x401260 <B::f()>:  push    %rbp
0x401261 <B::f()+1>:  mov     %rsp,%rbp
0x401264 <B::f()+4>:  mov     %rdi,-0x8(%rbp)
0x401268 <B::f()+8>:  mov     $0x7,%eax
0x40126d <B::f()+13>:  pop     %rbp
0x40126e <B::f()+14>:  retq
0x40126f:      nop
0x401270 <A::f()>:  push    %rbp
0x401271 <A::f()+1>:  mov     %rsp,%rbp
0x401274 <A::f()+4>:  mov     %rdi,-0x8(%rbp)
0x401278 <A::f()+8>:  mov     $0x5,%eax
0x40127d <A::f()+13>:  pop     %rbp
0x40127e <A::f()+14>:  retq
0x40127f:      nop
0x401280 <__libc_csu_init>:  endbr64
```

Inheritance

- In this case vptr is located at the start of an object
- May differ between compilers
- Compiler can optimize out vptr-call
- The pointer cast between base and derived class is free
- Its not true in case of multiple and/or virtual inheritance

Multiple inheritance

```
struct A {  
    int a;  
    virtual int f() { return 5; }  
};  
  
struct B {  
    int b;  
    virtual int g() { return 7; }  
};  
  
struct C: public A, public B {  
    int c;  
    C(int a, int b, int c): A(a), B(b), c(c) {}  
};  
  
C c{1, 2, 3};  
A *a = &c;  
B *b = &c;
```

Multiple inheritance

```
C c{1, 2, 3};  
A *a = &c;  
B *b = &c;
```

(gdb) x/60bx &b

0x7fffffff108:	0x28	0xe1	0xff	0xff	0xff	0x7f	0x00	0x00
0x7fffffff110:	0x18	0xe1	0xff	0xff	0xff	0x7f	0x00	0x00
0x7fffffff118:	0x18	0x20	0x40	0x00	0x00	0x00	0x00	0x00
0x7fffffff120:	0x01	0x00	0x00	0x00	0x00	0x00	0x00	0x00
0x7fffffff128:	0x30	0x20	0x40	0x00	0x00	0x00	0x00	0x00
0x7fffffff130:	0x02	0x00	0x00	0x00	0x03	0x00	0x00	0x00
0x7fffffff138:	0x00	0x00	0x00	0x00	0x00	0x00	0x00	0x00
0x7fffffff140:	0x00	0x00	0x00	0x00				

Bytes

- Types exist during compilation only, and any type in c++ is just a way to interpret bytes in memory.

```
struct Item {  
    int a;  
    int b;  
    int64_t c;  
    char payload[16];  
};  
  
char *data = ...; // assume it points somewhere we expect the same protocol  
                // a binary file for a example  
Item *item = reinterpret_cast<Item*>(data); // or (Item*)(data);
```

- This is **UB** (but nobody cares)

Bytes

- May lead to problems, but its very unlikely
- Don't forget about an alignment and little/big-endian though.
- The right way to do this (`bit_cast` since c++20)

```
Item item;  
memcpy(&item, data, sizeof(Item));
```

- The high 16 bits of pointers are not used on x86_64, so you can store some data there.
- The low bits of pointers can also be used due to the alignment of the underlying data.

Conclusion

- What we've omitted
 - virtual inheritance
 - pointers to class members (`class A` , what exactly is `&A::a`)
 - pointers to methods (what is `&A::method`)
- Additional materials
 - Inside the C++ Object Model By Stanley B. Lippman
 - https://en.wikipedia.org/wiki/Two's_complement
 - <https://graphics.stanford.edu/~seander/bithacks.html>