

WEEK 7

Perf

Perf

- perf stat
- perf record/report
- perf syscall

perf stat

- Allows you to measure various counters of different executables (`perf list`)
 - branches, branch-misses
 - cycles, instructions, context-switches
 - major-faults, minor-faults, faults
 - cache-misses, cache-references, L1-dcache-load-misses, L1-dcache-loads
- `perf stat -e <events>` to specify concrete events

Example

- Lets make a program, that takes sorted array and then processes multiple queries of "find lower_bound(x) for given x"
- Simple solution (binary search)

```
int Query(int x) const {  
    auto it = std::lower_bound(data_.begin(), data_.end(), x);  
    return it == data_.end() ? 0 : *it;  
}
```

- Assume the length of array is 10^8 and there are 10^7 queries
- With this array size the expected bottleneck is gonna be random memory jumps, that binary search does

perf stat

- perf stat by default doesn't show much:

```
> perf stat ./bin
Performance counter stats for './bin':

      8,401.37 msec task-clock                #    0.999 CPUs utilized
           76      context-switches         #    0.009 K/sec
            0      cpu-migrations            #    0.000 K/sec
        195,440    page-faults               #    0.023 M/sec
  26,730,393,359    cycles                    #    3.182 GHz
    44,953,734     stalled-cycles-frontend   #    0.17% frontend cycles idle
    85,475,051     stalled-cycles-backend    #    0.32% backend cycles idle
  11,219,240,562    instructions              #    0.42 insn per cycle
                                   #    0.01 stalled cycles per insn
    1,060,599,655    branches                 # 126.241 M/sec
            0      branch-misses            #    0.00% of all branches
```

- So lets try to check cache events

Example

```
>perf stat -e cache-references,cache-misses,L1-dcache-loads,L1-dcache-load-misses ./bin
Performance counter stats for './bin':
```

317,474,639	cache-references:u		
131,958,911	cache-misses:u	#	41.565 % of all cache refs
1,620,186,411	L1-dcache-loads:u		
165,314,470	L1-dcache-load-misses:u	#	10.20% of all L1-dcache hits

```
8.634930995 seconds time elapsed
```

```
8.375269000 seconds user
```

```
0.255977000 seconds sys
```

Optimization

- Lets split the initial array into blocks of size 10^4 and save maximum in each block separately

```
void InitIndex() {  
    block_size = sqrt(data_.size());  
    index_.reserve(data_.size() / block_size);  
    for (size_t i = 0; i < data_.size(); i += block_size) {  
        size_t end = std::min(i + block_size, data_.size());  
        index_.push_back(data_[end - 1]);  
    }  
}
```

Optimization

- To process a query, perform a binary search on `index_` first to find the corresponding block and then another binary search in this block

```
int Query(int x) const {
    auto it = std::lower_bound(index_.begin(), index_.end(), x);
    if (it == index_.end()) {
        return 0;
    }
    size_t block_num = it - index_.begin();
    size_t left = block_size * block_num;
    size_t right = std::min(data_.size(), left + block_size);
    return *std::lower_bound(data_.begin() + left, data_.begin() + right, x);
}
```


Optimization

```
>perf stat -e cache-references,cache-misses,L1-dcache-loads,L1-dcache-load-misses ./bin
Performance counter stats for './fast':
```

240,271,699	cache-references:u		
103,785,706	cache-misses:u	#	43.195 % of all cache refs
1,710,791,330	L1-dcache-loads:u		
126,991,538	L1-dcache-load-misses:u	#	7.42% of all L1-dcache hits

```
7.620999804 seconds time elapsed
```

```
7.374352000 seconds user
```

```
0.243945000 seconds sys
```

perf record

- `perf record` allows to profile given executable
- `perf report` builds a report based on the profiling's result
- Build with `-g -fno-omit-frame-pointer`
 - Sometimes `-O0` is needed if you want to see a proper call-graph
- `perf report` uses `cycles` by default, however you given specify an event with `-e`


Example

```
> clang++-15 -std=c++17 -O0 -g -fno-omit-frame-pointer bin.cpp -o bin
> perf record -g ./bin
took 15.9982
[ perf record: Woken up 3 times to write data ]
[ perf record: Captured and wrote 0.535 MB perf.data (5521 samples) ]
> perf report
```

```
Samples: 5K of event 'cycles:u', Event count (approx.): 2247109875
Children      Self  Command  Shared Object  Symbol
+ 100.00%    0.00% bin      libc-2.31.so   [.] __libc_start_main
- 99.45%    1.74% bin      bin            [.] main
- 97.71% main
- 70.92% Index::Query
- 69.90% std::lower_bound<__gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >, int>
- 68.95% std::__lower_bound<__gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >, int, __gnu_cxx::__ops::_Iter_less_val::operator()<__gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >, int>
+ 43.17% __gnu_cxx::__ops::_Iter_less_val::operator()<__gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >, int>
+ 10.06% std::advance<__gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >, long>
+ 1.24% __gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >::operator++
+ 22.86% std::uniform_int_distribution<int>::operator()<std::linear_congruential_engine<unsigned long, 48271ul, 0ul, 2147483647ul> >
+ 2.48% std::vector<int, std::allocator<int> >::vector
+ 0.81% __gnu_cxx::operator!=<int*, std::vector<int, std::allocator<int> > >
+ 1.74% __libc_start_main
+ 70.92%    0.22% bin      bin            [.] Index::Query
+ 69.90%    0.22% bin      bin            [.] std::lower_bound<__gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >, int, __gnu_cxx::__ops::_Iter_less_val::operator()<__gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >, int>
+ 68.95%   13.20% bin      bin            [.] std::__lower_bound<__gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >, int, __gnu_cxx::__ops::_Iter_less_val::operator()<__gnu_cxx::__normal_iterator<int const*, std::vector<int, std::allocator<int> > >, int>
```

Flamegraph

- Allows to view `perf report` in a more convenient and graphic way
- <https://github.com/brendangregg/FlameGraph>



perf syscall

- perf itself uses `perf_event_open` syscall

```
int syscall(SYS_perf_event_open, struct perf_event_attr *attr,  
pid_t pid, int cpu, int group_fd, unsigned long flags);
```

- Allows to analyze some block of code, not the entire program
- https://man7.org/linux/man-pages/man2/perf_event_open.2.html